

Probabilistic Programming Languages

Guillaume Baudart
Christine Tasson

MPRI 2021-2022

Overview

Probabilistic Programming Languages

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Bayesian Inference: learn parameters from data

- Latent parameter θ
- Observed data x_1, \dots, x_n

$$p(\theta \mid x_1, \dots, x_n) = \frac{p(\theta) p(x_1, \dots, x_n \mid \theta)}{p(x_1, \dots, x_n)} \quad (\text{Bayes' theorem})$$

posterior

$$\propto p(\theta) p(x_1, \dots, x_n \mid \theta) \quad (\text{Data are constants})$$

prior

likelihood



Thomas Bayes (1701-1761)

Probabilistic Programming Languages

General purpose programming languages extended with probabilistic constructs

- **sample**: draw a sample from a distribution
- **assume**, **factor**, **observe**: condition the model on inputs (e.g., observed data)
- **infer**: compute the posterior distribution of a model given the inputs

Multiple examples:

- Church, Anglican (lisp, clojure), 2008
- WebPPL (javascript), 2014
- Pyro/NumPyro (python), 2017/2019
- Gen (julia), 2018
- ProbZelus (Zelus), 2019
- ...

More and more, incorporating new ideas:

- New inference techniques, e.g., stochastic variational inference (SVI)
- Interaction with neural nets (deep probabilistic programming)

Inference in Practice

Overview

Rejection Sampling

```
module Rejection_sampling = struct
  exception Reject

  let sample d = Distribution.draw d
  let assume p = if not p then raise Reject

  let infer ?(n = 1000) model obs =
    let rec exec i = try model Prob obs with Reject → exec i in
    let values = Array.init n exec in
    Distribution.uniform_support ~values
end
```

Executing the model generates one sample

- **sample**: draw from a distribution
- **assume/observe**: hard conditioning, reject invalid samples
- Terminates with *n* valid samples

Importance Sampling

```
let coin prob data =  
  type prob = { id : int; scores : float array }  
  
  let sample _prob d = Distribution.draw d  
  let factor prob s = prob.scores.(prob.id) ← prob.scores.(prob.id) +. s  
  
  let infer ?(n = 1000) model obs =  
    let scores = Array.make n 0. in  
    let values = Array.mapi (fun i _ → model { id = i; scores } obs) scores in  
    Distribution.support ~values ~logits:scores  
end
```

Executing the model generates a pair (sample, weight)

- **sample**: draw from a distribution
- **factor/observe**: soft conditioning, assign a score
- Terminates with *n* pairs (sample, weight)

Particle Filter

basic.ml

```
module Particle_filter = struct
  include Importance_sampling

  let resample particles = ... (* resample an array of (particle, score) *)

  let factor s k prob = ...
    (* execute all the particle until the next factor statement *)
    (* resample given the score so far *)
    (* resume execution until the next factor statement *)
end
```

Inference algorithm : importance sampling, but...

- Add a resampling step at each **factor**
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution

Using CPS models, we can easily clone, stop, and resume particles the middle of an execution.

Coin


coin.ml

```
let _ =  
  let data = List.init 1000 (fun i → i mod 2) in  
  let dist = infer coin data in  
  plot dist
```


Rejection Sampling

Very (very) slow!

Importance Sampling



Particle Filter



Particle impoverishment

HMM

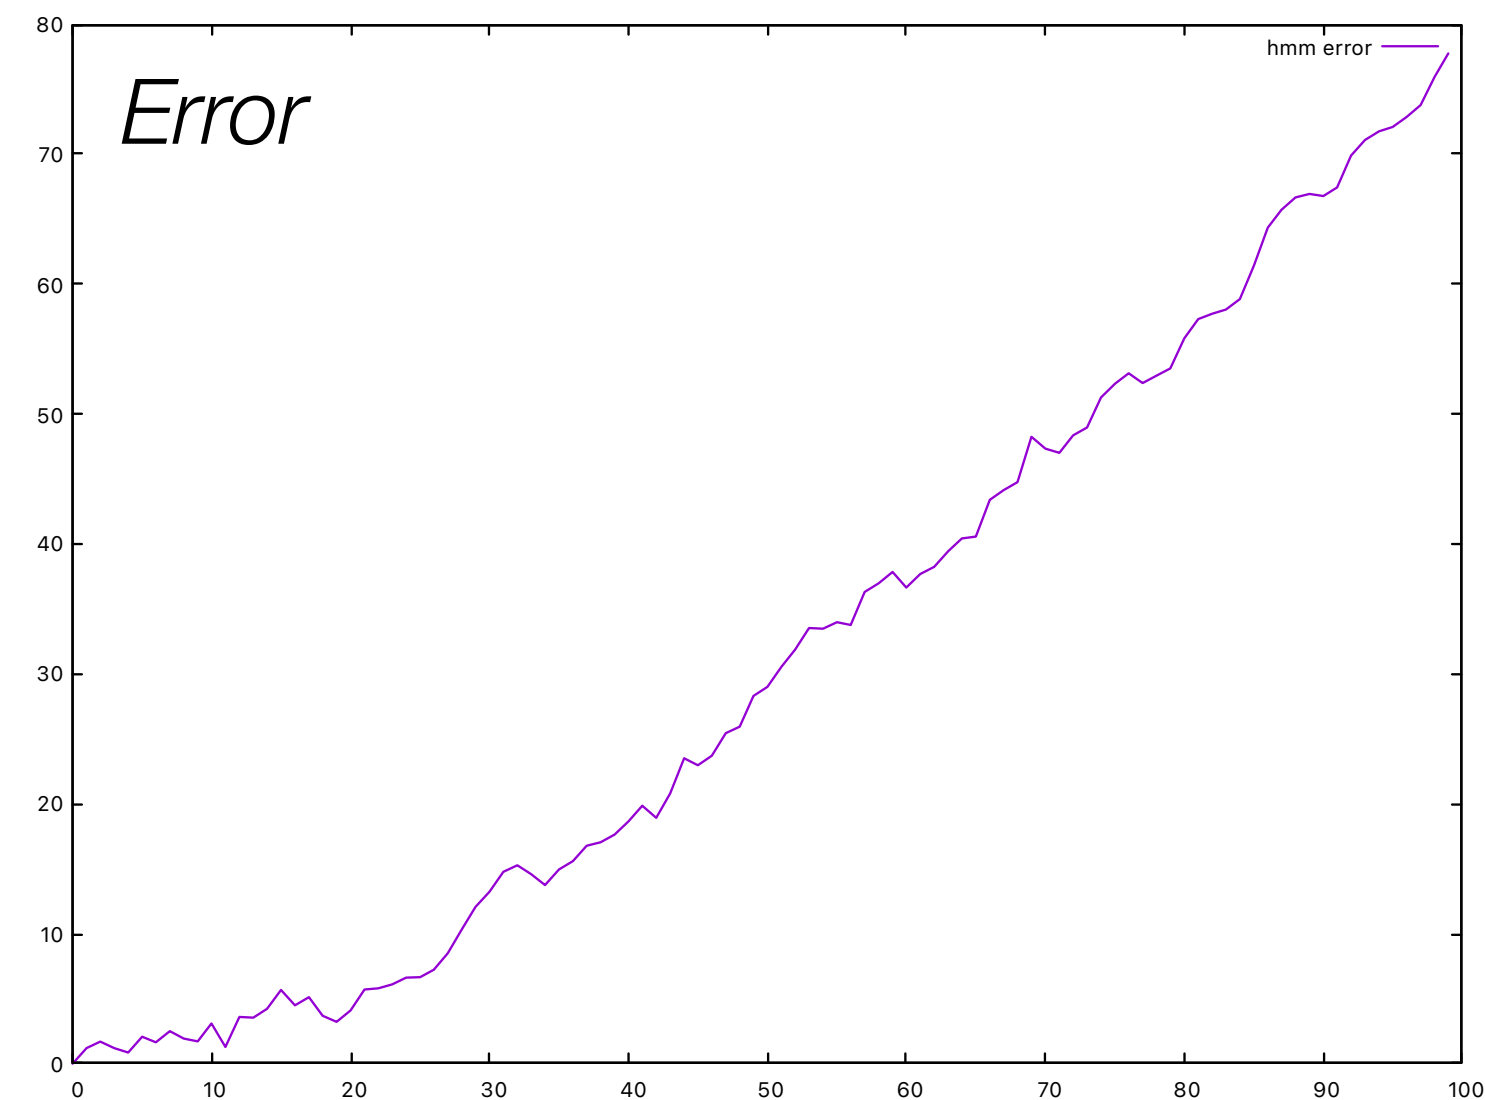
hmm.ml

```
let _ =  
  let data = Owl.Arr.linspace 0. 100. 100 ▷ Owl.Arr.to_array ▷ Array.to_list in  
  let dist = Distribution.split_list (infer hmm data) in  
  plot (error (List.rev dist) data)
```

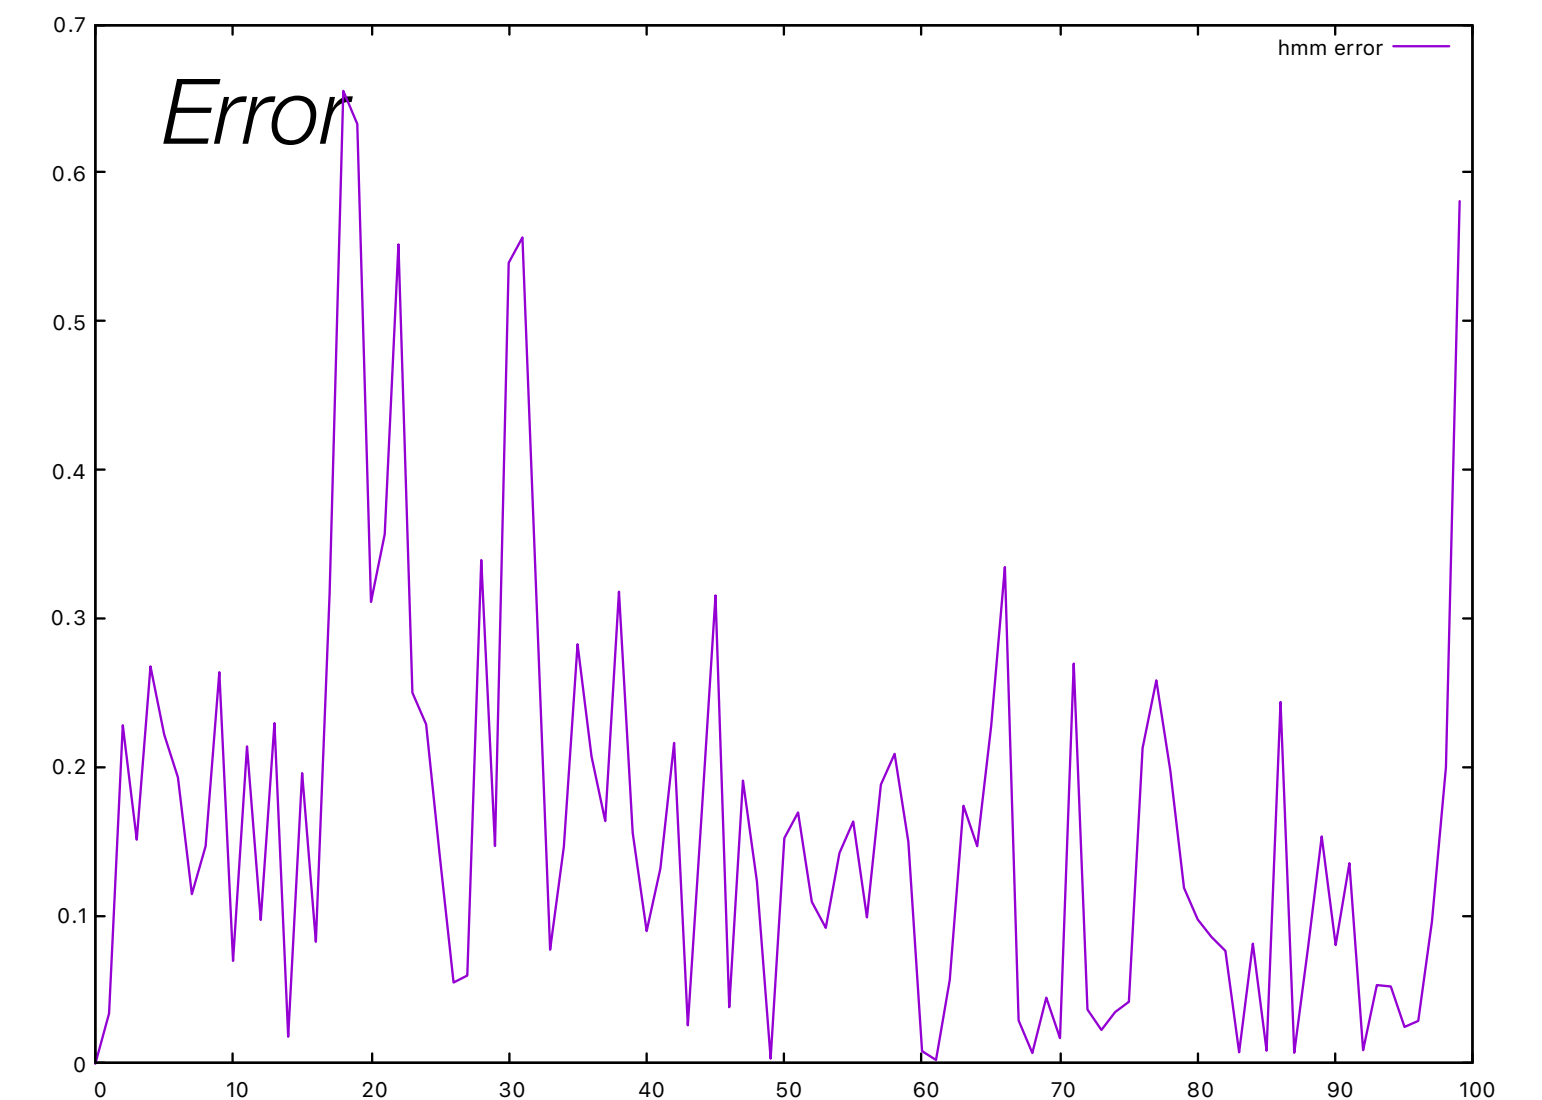
Rejection Sampling

Never terminate!

Importance Sampling



Particle Filter



Denotational Semantics

Overview

Deterministic vs. Probabilistic

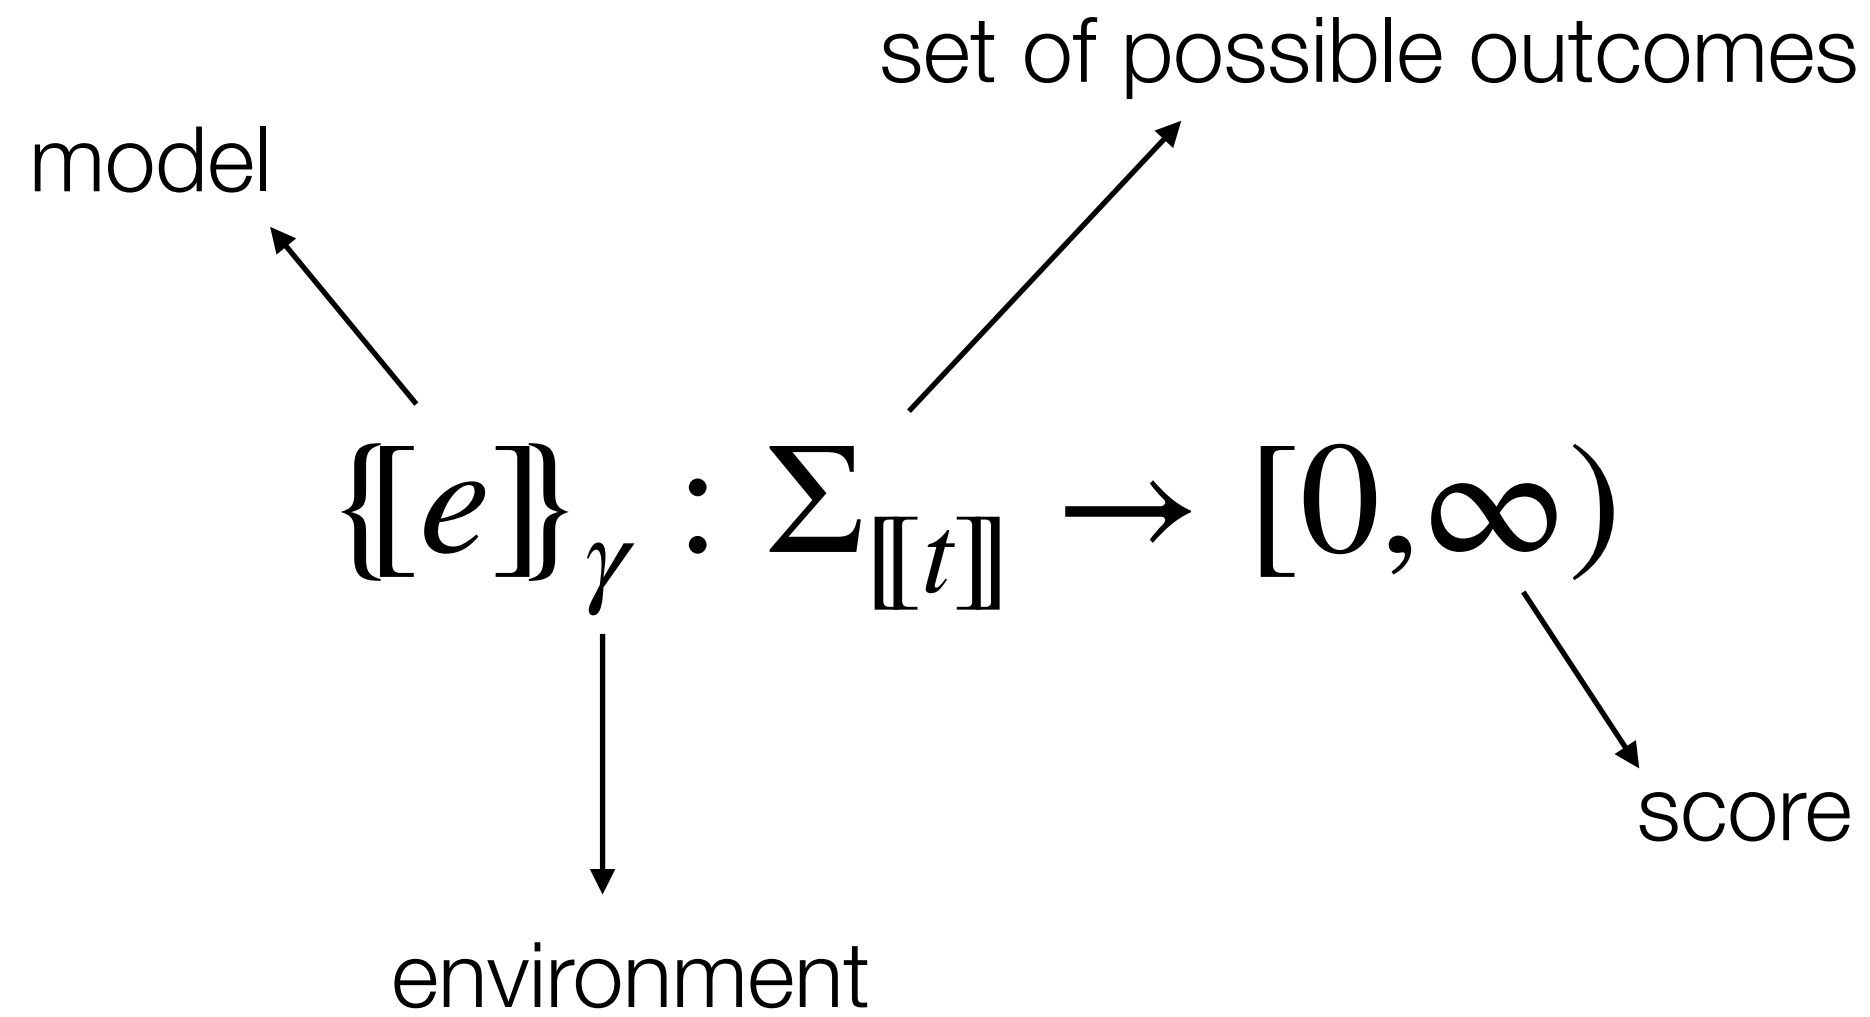
Deterministic semantics $G \vdash^{\text{D}} e : t$

- Classic denotational semantics
- Given an environment $\gamma : \Gamma$, $\llbracket e \rrbracket_{\gamma}$ returns a value of type t
- $\llbracket e \rrbracket : \Gamma \rightarrow t$

Probabilistic semantics $G \vdash^{\text{P}} e : t$

- Expressions are interpreted as kernels
- Given an environment $\gamma : \Gamma$, $\{\!\{ e \}\!\}_{\gamma}$ is a measure on values of type t
- $\{\!\{ e \}\!\} : \Gamma \times \Sigma_{\llbracket t \rrbracket} \rightarrow [0, \infty)$

Probabilistic Programming



Unnormalized measure

$$\llbracket \text{infer}(e) \rrbracket_\gamma = \frac{\llbracket e \rrbracket_\gamma}{\llbracket e \rrbracket_\gamma(\llbracket \text{typeOf}(e) \rrbracket)}$$

Distribution

normalize over all possible values

Probabilistic Semantics

$$\llbracket e \rrbracket_\gamma = \lambda U. \delta_{\llbracket e \rrbracket_\gamma}(U) \text{ if } \text{kindOf}(e) = \text{D}$$

$$\llbracket \text{let } p = e_1 \text{ in } e_2 \rrbracket_\gamma = \lambda U. \int_{\llbracket \text{typeOf}(e_1) \rrbracket} \llbracket e_1 \rrbracket_\gamma(du) \llbracket e_2 \rrbracket_{\gamma+[u/p]}(U)$$

$$\llbracket \text{sample}(e) \rrbracket_\gamma = \lambda U. \llbracket e \rrbracket_\gamma(U)$$

$$\llbracket \text{factor}(e) \rrbracket_\gamma = \lambda U. \exp(\llbracket e \rrbracket_\gamma) * \delta_{()}(U)$$

Example : Beta

my_gaussian.ml

```
let my_beta a b =  
  let x = sample (uniform ~a:0 ~b:1) in  
  let () = observe (beta ~a ~b) x in  
  x
```

$$\begin{aligned}\llbracket \text{my_beta } a \text{ } b \rrbracket_-(U) &= \int_0^1 \llbracket \text{sample (uniform } 0 \text{ } 1) \rrbracket_{\{a/a, b/b\}}(dx) \\ &\quad \int_{()} \llbracket \text{observe (beta } a \text{ } b) \text{ } x \rrbracket_{\{a/a, b/b, x/x\}}(du) \llbracket x \rrbracket_{\{a/a, b/b, x/x, u/()\}}(U) \\ &= \int_0^1 \text{Uniform}(dx) \text{Beta}_{\text{pdf}}(x) \delta_x(U) \\ &= \int_U \text{Beta}(a, b)_{\text{pdf}}(x) dx \\ &= \text{Beta}(a, b)(U)\end{aligned}$$

Operational Semantics

Overview

Sampler

Probabilistic semantics $G \vdash^P e : t$

- Expressions are interpreted as samplers
- Given an environment $\gamma : \Gamma$, and a weight $w \in [0, \infty)$, $\llbracket e \rrbracket_{\gamma, w} = v, w'$ returns a pair (value, new score)
- $\llbracket e \rrbracket : \Gamma \times [0, \infty) \rightarrow [0, \infty)$

$$\llbracket c \rrbracket_{\gamma, w} = c, w$$

$$\llbracket x \rrbracket_{\gamma, w} = \gamma(x), w$$

$$\llbracket \text{sample}(e) \rrbracket_{\gamma, w} = \text{let } d, w' = \llbracket e \rrbracket_{\gamma, w} \text{ in draw}(d), w'$$

$$\llbracket \text{factor}(e) \rrbracket_{\gamma, w} = \text{let } s, w' = \llbracket e \rrbracket_{\gamma, w} \text{ in } (), s * w'$$

Importance Sampling

Inference algorithm

- Run a set of n independent executions
- **sample**: draw a sample from a distribution
- **factor**: associate a score to the current execution
- Gather output values and score to approximate the posterior distribution

$$\llbracket \text{infer}(e) \rrbracket_{\gamma} = \underbrace{\text{let } [(v_i, w'_i) = \{e\}_{\gamma, w_i}]_{1 \leq i \leq N}}_{\text{particles}} \text{ in } \underbrace{\lambda U. \sum_{1 \leq i \leq N} \overline{w'_i} * \delta_{v_i}(U)}_{\text{distribution}}$$

Normalized weights: $\overline{w}_i = \frac{w_i}{\sum_{i=1}^n w_i}$

Reactive Probabilistic Programming

Overview

Reactive Probabilistic Systems

Synchronous data-flow languages and block diagrams

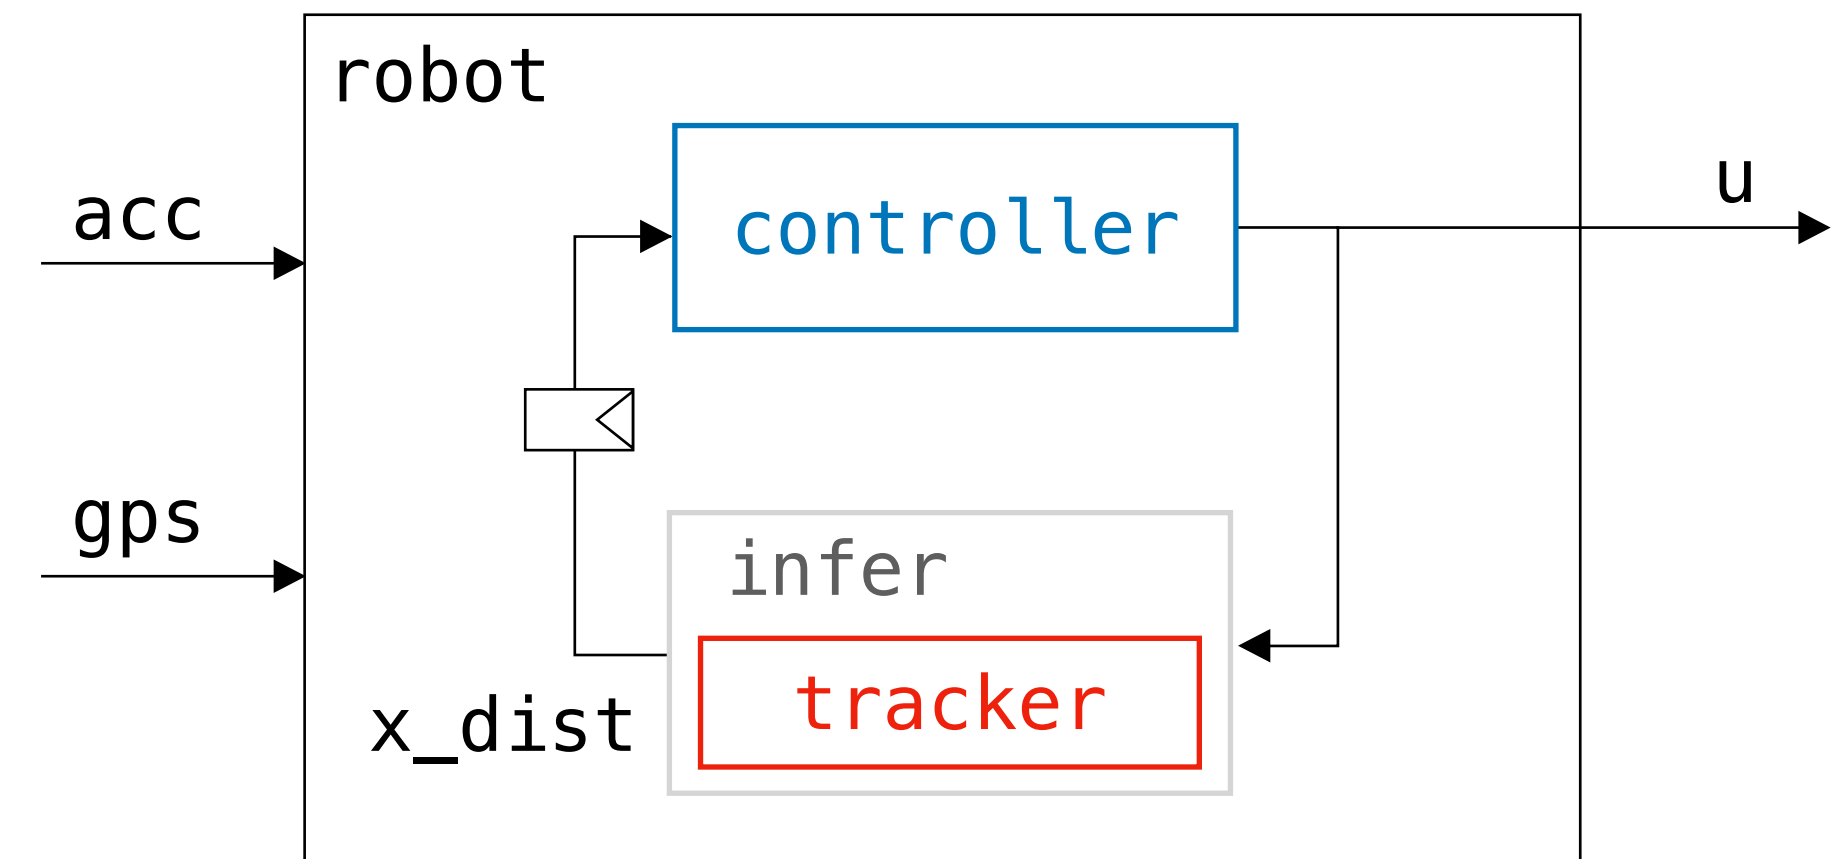
- Signal: stream of values
- System: stream processor

ProbZelus: add support to deal with uncertainty

- Extend a synchronous language
- Parallel composition: deterministic/probabilistic
- Inference-in-the-loop
- Streaming inference

Probabilistic programming with streams

State: $x_dist: (position \times velocity \times acceleration) dist$



```
let node robot (acc, gps) = u where
  rec u = controller (x0_dist → pre x_dist)
  and x_dist = infer tracker (u, acc, gps)
```


Reactive Probabilistic Programming

Language design and implementation

- Type system deterministic / probabilistic
- Measure-based formal semantics
- Compilation to transition functions

Streaming inference

- Sequential Monte-Carlo method for reactive models
- Semi-symbolic delayed sampling with bounded memory



Evaluation

- Multiple benchmarks from simple to advanced
- Inference schemes comparison: speed vs. accuracy



<https://github.com/IBM/probzelus>